


Multimodal attention-fusion model: training curves

Loss



Validation accuracy

